

# Security Project: Fuzzing and Symbolically Executing Cryptographic Implementations

Heman Gandhi Rutgers University

## Abstract

Many programs assume that they implement the common secure encryptions protocols such as RSA and AES. However, it is difficult for implementors to test their beliefs: many implementations of the algorithms can be faulty. This can have untold consequences for users of the implementations.

Fuzzing is a great tool for finding the older exploits that used to wreak havoc: fuzzers can find buffer overflows and various other attacks dotting large low-level code bases.

Fuzzing to solve the security holes in encryption programs is difficult, however. The issues occur over a small fraction of the cases, unlike buffer overflows. To make the two approaches more compatible, symbolic execution might be able to constraint the test-case generation to produce tests that are more likely to trip up cryptographic applications. Another source of difficulty is the reliance on randomness.

This work hopes to detail some approaches and difficulties in using fuzzing to solve the aforementioned security holes.

This implementation is available at [this GitHub](#).

## 1 Introduction

Fuzzing can be implemented in various ways though the canonical choice is AFL [1]. As the site suggests, fuzzing is a strong way to find bugs in many low-level programs.

In addition to buffer overflows and side channel attacks, in the last few decades, the security of cryptography has been called into question. [2] However, to test such implementations, such as the entropy of random number generators, fuzzing does not suffice.

Here, a middle ground is explored and simpler alternatives are analysed. While initially thought to be the simplest direct approach, fuzzing Java informed by symbolic execution turns out to be overly complicated. It may be simpler to statically generate the test cases. This is discussed in the later sections: section 2 details the background on fuzzing, symbolic execution, and how their union would test the security of cryptographic code; section 3 details the experimental set up with JFP, SymExe, and badger; section 4 discusses an alternate approach that may be simpler for initially using symbolic execution; section 5 concludes.

## 2 Related Works

There are a two classes of related works: works on fuzzing, and works on breaking cryptographic protocols. Finally, there is also their meet: works on fuzzing against cryptographic code.

### 2.1 Fuzzing and Symbolic Execution

Fuzzing is a simple technique at its core: it is to generate test cases and run the program until the program fails. This, however, is difficult to reason about for various reasons. The generation of inputs is the biggest concern: it is very difficult to generate test cases that cover some black-boxed executable. Various approaches are discussed here.

First, there is the classical AFL approach. [1] is the conventional fuzzer used by many to test low-level libraries. On its site is a large trophy case of bugs found by the genetic algorithm at the heart of this fuzzer. The algorithm tracks coverage by labelling basic blocks and then using the xor of the blocks to

maintain the edge between them. The collection of edges forms a fairly unique representation of an execution trace. If an execution trace does not lead to a new behavior, the test case is not further explored. Otherwise, if new behavior is found, the genetic algorithm is used to further explore the input.

Another approach is to compare two runs of the program at a higher level. This approach is useful for finding side-channels. Fuzzers for this have been shown to be effective [3] [4]. Here, [4] concludes that various Go, Python, and C++ libraries do not insure users against poor parameter choices and sometimes do not fully adhere to standards.

The two approaches above are powerful, however, here, symbolic execution is explored. Symbolic execution would allow the fuzzer to powerfully reason about number theoretic invariants in the various cryptographic algorithms implemented. This has not been done directly. Indirectly, much tooling is provided. In particular, [5] runs a trinity of processes to fuzz Java code. There is a fuzzer and a symbolic executor. These are joined in badger itself that uses the symbolic executor to inform the fuzzer to better test and execute Java. Simply: the symbolic executor finds program paths and the conditions that would cause the path to be executed. This is put into a trie and the leaf nodes are added to a queue. A leaf is dequeued, turned into a test case (by any SAT-solver) and that test case is passed to the fuzzer.

Hence, the fuzzing is an active field with various applications to testing cryptographic code.

## 2.2 Cryptographic Algorithms

Cryptographic algorithms power the modern internet. Many implementations of varied algorithms are used to secure many layers of the OSI stack: routers sign off their identity, servers do the same with SSL and TLS, and at the message and application layer, HTTPS encrypts the messages.

RSA is a common algorithm that is widely popular. For security, however, it depends on large primes. Furthermore, in the internet, it needs that all the primes found are independently found. For instance, if a server uses  $P$  and  $Q$  and another server uses  $Q$  and  $R$ , and all three, on their own are secure,

there is still a vulnerability. Namely, if an attacker knows this about the server, they can quickly compute  $gcd(PQ, QR)$  ( $PQ$  and  $QR$  being public as a part of RSA). This would give the attacker  $Q$  from where finding  $P$  and  $R$  is a matter of division. This issue can be endemic to a server, if the random number generator is poor enough. This is a widespread issue, as found in [2]. A fuzzer could be informed by their work and resources to ensure that primes used are most likely not vulnerable to this  $gcd$  attack.

This is suggestive of various invariants and issues that may occur in implementations. For this work, only RSA and the entropy of random number generators are considered. These should be a good starting point in a lengthier investigation of cryptographic algorithms.

## 2.3 The Union of the Approaches

As suggested above, differential fuzzing has proven useful for revealing side-channels and even found bugs in cryptography libraries [4]. That said, there is a bigger discussion of symbolic execution used to verify cryptography.

In fact, network protocols can be tested through symbolic execution [6]. Here, it is also possible to statically detect the use of bad pseudo-random number generators.

This work is an attempt to extend this fledgling field. Here, it is established that badger’s architecture [5], is too complex for quickly investigating whether symbolic execution is feasible – a more foundational approach would be more suitable.

## 3 The Program

Here, the structure and the necessary changes to badger are documented. Due to the complexity of the changes two things were impossible: searching using SymExe for the correct invariants, and foundationally building up a set of searches and paths that could find problems with cryptography implementations.

As an initial test, a publicly provided RSA was used [7]. It is documented that if  $message \geq modulus$  or  $gcd(\phi, pq) \neq 1$ , the encryption fails. The

hope was to extend badger so that these two bugs and possibly a bad pseudo-random number generator could be detected.

### 3.1 The Architecture of Badger

Badger is a behemoth. It runs a trinity of rather large Java libraries. A dependency graph is likely the best description:

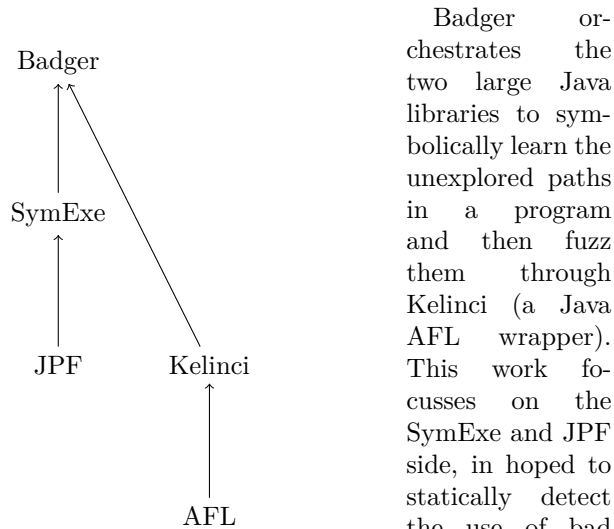


Figure 1: The badger trinity and dependencies

searches to explore invariants in cryptographic code.

Below, the JPF property added is mentioned with a few potential extensions. Then the infeasibility of extending SymExe is mentioned. (Note that SymExe is a part of badger, not a library.)

### 3.2 Extending JPF

The first work was to extend JPF to ensure that all random generators used were secure. This was not done too thoroughly for a first implementation: JPF provides properties that are checked at every step of their VM's execution. Using this, another property was created to check that `Math.random` and `java.util.Random` were not used (both are known to be insecure). The property queries the call stack

and ensures that no methods in `java.util.Random` are not called by checking that the class the method executing is in is not `java.util.Random`. Similarly, the property asks the call stack if it has entered `Math.random`.

However, JPF does not provide enough of a native interface to emulate `java.security.SecureRandom`. The file creation fails since it is impossible to emulate `/dev/random` purely symbolically. This would have to be added.

### 3.3 SymExe Searches

SymExe performs searches by expanding a trie of potential execution paths. Nodes that are high and are at branch points are, by the default metric, considered the most worth exploring.

Depending on the cryptographic algorithm in mind, SymExe could be extended to also create more paths based on the numbers processed: the numbers sharing common factors would be one of the potential paths, for instance. This would cover one of the bugs mentioned in the source code [7].

However, this approach does not generalise as most of the paths would depend on random numbers that were generated elsewhere. Hence, one would have to consider these cases and accurately model randomness.

### 3.4 Takeaways

The biggest lesson is perhaps a change in approach. Java is rather clunky by design and the dependency tree is quite deep for initial experimentation. A couple approaches would be a more feasible starting point:

- Instrumenting a C AST for AFL to fuzz cryptographic code.
- Using a lighter symbolic executor to merely generate the path conditions (as a proof-of-concept).
- Using more broken implementations or simpler sub-parts of algorithms for testing.

Either of these can also be pursued in a simpler language like C, where the security of the randomness generator cannot be obfuscated by sub-classing or other complicated structures.

Additionally, the approach has limited visibility into the workings of `BigInteger` which is used in Java to compute with large numbers. C would be a better alternative for this as these implementations would be easier to target.

Finally, an important thought is that randomness is very difficult to emulate. Without much external information, it is difficult to guess the entropy of random numbers. For symbolic execution, it may be necessary to treat a random number as an input fully controllable by the symbolic executor, but it would be more informative to propagate probabilities and similar details (see the potential future work below).

## 4 Future Work

There are many ways – in addition to the “short-term” takeaways above that could apply to cryptographic algorithms. These would be large projects predicated on the success of the symbolic execution fuzzing approach. These also bring up questions about cryptographic algorithms and how general any such system can be.

The two systems conceived of below are more static analysis extensions for fuzzing cryptographic algorithms.

### 4.1 Generally Specifying Algorithms

Technically, an implementation should be a specification of an algorithm. Of course, the reason such work is pursued, is that implementations often are not completely correct.

However, in order to test an implementation of a number theory algorithm, some external information might be necessary. Automatically inferring that a user may be able to solve for the public key, per say, would not be productive. Hence, some system of annotations should be considered so that a static analysis would know how an attacker would be approaching a particular implementation.

This can be done through a DSL that generates the input. Such an approach has proven to be effective in [8] for finding worst-case run times of various algorithms. A similar idea may extend to fuzzing cryptographic algorithms are a DSL can specify the conditions that would favor an attacker and these can be fed to the program to see if it is vulnerable. Naturally, such a DSL would also specify what parameters are publicly knowable and what parameters are intended to be kept secret.

### 4.2 Symbolic Randomness and Propagating Entropy

Another part of statically determining the security of an implementation is to have the entropy of random number generators known to the static analysis tool. As the random numbers are used in an algorithm, the entropy can be tracked – through static analysis or symbolic execution (if not both). This would allow tools to be aware of how random a given generator is.

This could tie with the above where the DSL may specify a minimum randomness in the output it seeks. Then executions can fuzz by concretely executing the program and also symbolically track the entropy. Either failing would then be considered a failure of the code.

Such an approach would be supplemented by probabilistic programming. This is detailed in [9] where various applications are discussed. Generally, such programming is used to discuss distributions and contain them as first class objects – so that sampling and reasoning about said sampling are made simpler. These languages are a great model for propagation of entropy and the other arguments above.

## 5 References

### References

- [1] Technical “whitepaper” for afl-fuzz.
- [2] Nadia Heninger, Zakier Durumeric, Eric Wustrow, J Alex Halderman. Mining Your Ps and Qs: De-

- tection of Widespread Weak Keys in Network Devices. *USENIX*, 2012.
- [3] Slurin Nilizadeh, Yannic Noller, Corina S. Pasareanu. DifFuzz: Differential Fuzzing for Side-Channel Analysis. *arXiv*.
  - [4] JP Aumasson, Yolan Romain. Automated Testing of Crypto Software Using Differential Fuzzing. *Blackhat*.
  - [5] Yannic Noller, Rody Kersten, Corina S. Pasareanu. Badger: Complexity Analysis with Fuzzing and SymbolicExecution. *ACM SIGSOFT International Symposium on Software Testing and Analysis*, 2018.
  - [6] Mathy Vanhoef, Frank Piessens. Symbolic Execution of Security Protocol Implementations: Handling Cryptographic Primitives. *arXiv*.
  - [7] Robert Sedgewick, Kevin Wayne. RSA.java.
  - [8] Jiayi Wei, Jia Chen, Yu Feng, Kostas Ferles, Isil Dilling. Singularity: Pattern Fuzzing for Worst Case Complexity. *ESEC/FSE*, 2018.
  - [9] Andrew D Gordon, Thomas A Henzinger, Aditya V Nori, Sriram K Rajamani. Probabilistic Programming. 2014.